Learning Deep Representations of Fine-Grained Visual Descriptions

Liangjie Cao

27 May 2018

1. Deep Structured Joint Embedding

Intuitively, the maximize the compatibility between a description and its matching image, and minimize compatibility with images from other classes. The objective is: given data S = {(v_n, t_n, y_n), n=1,..., N} containing visual information $v \in \vartheta$, text descriptions $t \in \tau$ and class labels $y \in \Upsilon$. They seek to learn functions f_v: $\vartheta \to \Upsilon$ and f_t:$\tau \to \Upsilon$ minimize the empirical risk.

$$\frac{1}{N} \sum \triangle (y_n, f_n(v_n)) + \triangle(y_n, f_n(t_n))$$

where $\triangle$: $\Upsilon x \Upsilon \to \Re$ is the 0-1 loss. Note that N is the number of image and text pairs in the training set, and so a given image can have multiple corresponding captions. In practice they have many visual descriptions and many images per class. During training, in each mini-batch we ?rst sample an image from each class, and then sample one of its ten corresponding captions. Since their text encoder models are all differentiable, they backpropagate (sub)-gradients through all text network parameters for end-to-end training. For the image encoder, they keep the network weights ?xed to the original GoogLeNet.


Figure 1. **Their model learns a scoring function between images and text descriptions.**

2. Text encoder models

In this section the authors describe the deep neural language models that we use for representing ?ne-grained visual descriptions. They compare the performance on zero-shot prediction tasks in Section 3. Text-based convolutional neural networks were studied in depth in [3] for the task of document classi?cation. The text-based CNN can be viewed as a standard CNN for images, except that the image width is 1 pixel and the number of channels is equal to the alphabet size. The 2D convolution and spatial max-pooling are replaced by temporal (1D) convolution and temporal max-pooling. After each convolution layer, they use recti?ed linear activation unit (ReLU), the overall network is constructed using convolution, pooling and thresholding activation function layers, followed by ful-

lyconnected layers to project onto the embedding space. The text embedding function is thus simply $\psi(t)$=CNN(t); the ?nal hidden layer of the CN-N. Figure 2 illustrates the convolutional-recurrent approach. The ?nal encoded feature is the average hidden unit activation over the sequence. The resulting scoring function can be viewed as a linear accumulation of evidence for compatibility with a query image (illustrated in Figure 1). It is also a linearized version of attention over the text sequence. This has the advantage that at test time for classi?cation or retrieval, one can use the averaged hidden units as a feature, but for diagnostic purposes one can backtrace the score computation to each time step of text processing. They also evaluate a baseline that represents descriptions using unsupervised word embeddings learned by word2vec [2]. Previous works on visual-semantic embedding have directly used the word embedings of target classes for zero-shot learning tasks. However, in their case we have access to many visual descriptions, and they would like to extract vector representations of them in real time; *i.e.* without re-running word2vec training. A very simple way to do this is to average the word embeddings of each word in the visual description. Although this loses the structure of the sentence, this nevertheless yields a strong baseline and in practice performs similarly to bag of words. The CUB dataset also has per-image attributes, but they found that using these does not improve performance compared to using a single averaged attribute vector per class.


Figure 2. **Their proposed convolutional-recurrent net**

3. Experimental results

In this section they describe our experiments on the Caltech-UCSD Birds dataset (CUB) and Oxford Flowers102 (Flowers) dataset. CUB contains 11,788 bird images from 200 different categories. Flowers contains 8189 ?ower images from 102 different categories. Following [1], the images in CUB are split into 100 training, 50 validation, and 50 disjoint test categories. The CNN input size (sequence length) was set to 30 for word-level and 201 for character-level models; longer text inputs are cut off at this point and shorter ones are zeropadded. All text embeddings used a 1024-dimensional embedding layer to match the size of the image embedding. They kept the image encoder ?xed, and used RMSprop with base learning rate 0.0007 and minibatch size 40. I will continue learning in the following days.

References

[1] Zeynep Akata, Florent Perronnin, Zaid Harchaoui, and Cordelia Schmid. Label-embedding for image classification. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 38(7):1425–1438, 2016.

2

[2] Tomas Mikolov, Ilya Sutskever, Kai Chen, Gregory S Corrado, and Jeffrey Dean. Distributed representations of words and phrases and their compositionality. *neural information processing systems*, pages 3111–3119, 2013.

[3] Xiang Zhang, Junbo Jake Zhao, and Yann Lecun. Character-level convolutional networks for text classification. *neural information processing systems*, pages 649–657, 2015.